

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [2]: df = pd.read_csv('train.csv')
df.head()
```

Out[2]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	LP001002	Male	No	0	Graduate	No	5849	
1	LP001003	Male	Yes	1	Graduate	No	4583	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	
4	LP001008	Male	No	0	Graduate	No	6000	

```
In [3]: df.info
```

Out[3]: <bound method DataFrame.info of

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	No
1	LP001003	Male	Yes	1	Graduate	No	No
2	LP001005	Male	Yes	0	Graduate	Yes	Yes
3	LP001006	Male	Yes	0	Not Graduate	No	No
4	LP001008	Male	No	0	Graduate	No	No
5	LP001011	Male	Yes	2	Graduate	Yes	Yes
6	LP001013	Male	Yes	0	Not Graduate	No	No
7	LP001014	Male	Yes	3+	Graduate	No	No
8	LP001018	Male	Yes	2	Graduate	No	No
9	LP001020	Male	Yes	1	Graduate	No	No
10	LP001024	Male	Yes	2	Graduate	No	No
11	LP001027	Male	Yes	2	Graduate	NaN	NaN
12	LP001028	Male	Yes	2	Graduate	No	No
13	LP001029	Male	No	0	Graduate	No	No
14	LP001030	Male	Yes	2	Graduate	No	No
15	LP001032	Male	No	0	Graduate	No	No
16	LP001034	Male	No	1	Not Graduate	No	No
17	LP001036	Female	No	0	Graduate	No	No

In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
Loan_ID                614 non-null object
Gender                 601 non-null object
Married                611 non-null object
Dependents              599 non-null object
Education              614 non-null object
Self_Employed          582 non-null object
ApplicantIncome        614 non-null int64
CoapplicantIncome      614 non-null float64
LoanAmount             592 non-null float64
Loan_Amount_Term       600 non-null float64
Credit_History         564 non-null float64
Property_Area          614 non-null object
Loan_Status            614 non-null object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.4+ KB
```

In [5]: `df.describe()`

Out[5]:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
<b>count</b>	614.000000	614.000000	592.000000	600.000000	564.000000
<b>mean</b>	5403.459283	1621.245798	146.412162	342.000000	0.842199
<b>std</b>	6109.041673	2926.248369	85.587325	65.12041	0.364878
<b>min</b>	150.000000	0.000000	9.000000	12.000000	0.000000
<b>25%</b>	2877.500000	0.000000	100.000000	360.000000	1.000000
<b>50%</b>	3812.500000	1188.500000	128.000000	360.000000	1.000000
<b>75%</b>	5795.000000	2297.250000	168.000000	360.000000	1.000000
<b>max</b>	81000.000000	41667.000000	700.000000	480.000000	1.000000

In [6]: `df['Property_Area'].value_counts()`


Out[6]:

Semiurban	233
Urban	202
Rural	179

Name: Property\_Area, dtype: int64


```
In [7]: df['ApplicantIncome'].hist(bins=50)
```

```
Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x21b65cc76d8>
```




```
In [8]: df.boxplot(column='ApplicantIncome')
```

```
Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x21b65ddca90>
```




```
In [9]: df['LoanAmount'].hist(bins=50)
```

```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x21b660b5e80>
```



```
In [10]: df.boxplot(column='LoanAmount')
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x21b66175128>
```



```
In [11]: temp1 = df['Credit_History'].value_counts(ascending=True)
temp2 = df.pivot_table(values='Loan_Status',index=['Credit_History'],aggfunc=lambda x: x.mode()[0])
print ('Frequency Table for Credit History:')
print (temp1)

print ('\nProbability of getting loan for each Credit History class:')
print (temp2)
```

Frequency Table for Credit History:

0.0 89

1.0 475

Name: Credit\_History, dtype: int64

Probability of getting loan for each Credit History class:

Credit_History	Loan_Status
0.0	0.078652
1.0	0.795789

```

In [13]: fig = plt.figure(figsize=(8,4))

ax1 = fig.add_subplot(121)
ax1.set_xlabel('Credit_History')
ax1.set_ylabel('Count of Applicants')
ax1.set_title("Applicants by Credit_History")
temp1.plot(kind='bar')


ax2 = fig.add_subplot(122)
temp2.plot(kind = 'bar')
ax2.set_xlabel('Credit_History')
ax2.set_ylabel('Probability of getting loan')
ax2.set_title("Probability of getting loan by credit history")

```

```

Out[13]: Text(0.5,1,'Probability of getting loan by credit history')

```



```
In [14]: df.apply(lambda x: sum(x.isnull()),axis=0)
```

```
Out[14]: Loan_ID          0
Gender          13
Married         3
Dependents      15
Education       0
Self_Employed   32
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount      22
Loan_Amount_Term 14
Credit_History  50
Property_Area   0
Loan_Status     0
dtype: int64
```

```
In [15]: df['LoanAmount'].fillna(df['LoanAmount'].mean(), inplace=True)
```

```
In [16]: df['Self_Employed'].fillna('No',inplace=True)
```

```
In [17]: table = df.pivot_table(values='LoanAmount', index='Self_Employed', columns='Education')
# Define function to return value of this pivot_table
def fage(x):
    return table.loc[x['Self_Employed'],x['Education']]
# Replace missing values
df['LoanAmount'].fillna(df[df['LoanAmount'].isnull()].apply(fage, axis=1), inplace=True)
```

```
-----
TypeError                                Traceback (most recent call last)
<ipython-input-17-7b68a5f9512f> in <module>()
      4 return table.loc[x['Self_Employed'],x['Education']]
      5 # Replace missing values
----> 6 df['LoanAmount'].fillna(df[df['LoanAmount'].isnull()].apply(fage, axis=
1), inplace=True)

C:\anaconda\lib\site-packages\pandas\core\series.py in fillna(self, value, method, axis, inplace, limit, downcast, **kwargs)
    3420             axis=axis, inplace=inplace,
    3421             limit=limit, downcast=downcast,
t,
-> 3422             **kwargs)
    3423
    3424 @Appender(generic._shared_docs['replace'] % _shared_doc_kwargs)


C:\anaconda\lib\site-packages\pandas\core\generic.py in fillna(self, value, method, axis, inplace, limit, downcast)
    5398         raise TypeError('"value" parameter must be a scalar, dict or Series, but you passed a '
r, dict '
    5399                             '{0}'.format(type(value).__name__)
-> 5400     )
    5401
    5402         new_data = self._data.fillna(value=value, limit=limit,

TypeError: "value" parameter must be a scalar, dict or Series, but you passed a
"DataFrame"
```




```
In [18]: df['LoanAmount_log'] = np.log(df['LoanAmount'])  
df['LoanAmount_log'].hist(bins=20)
```

Out[18]: <matplotlib.axes.\_subplots.AxesSubplot at 0x21b67628940>



```
In [19]: df['TotalIncome'] = df['ApplicantIncome'] + df['CoapplicantIncome']  
df['TotalIncome_log'] = np.log(df['TotalIncome'])  
df['LoanAmount_log'].hist(bins=20)
```

Out[19]: <matplotlib.axes.\_subplots.AxesSubplot at 0x21b676f77f0>



```
In [20]: df['Gender'].fillna(df['Gender'].mode()[0], inplace=True)  
df['Married'].fillna(df['Married'].mode()[0], inplace=True)  
df['Dependents'].fillna(df['Dependents'].mode()[0], inplace=True)  
df['Loan_Amount_Term'].fillna(df['Loan_Amount_Term'].mode()[0], inplace=True)  
df['Credit_History'].fillna(df['Credit_History'].mode()[0], inplace=True)
```

```
In [21]: from sklearn.preprocessing import LabelEncoder
var_mod = ['Gender', 'Married', 'Dependents', 'Education', 'Self_Employed', 'Property_Area']
le = LabelEncoder()
for i in var_mod:
    df[i] = le.fit_transform(df[i])
df.dtypes
```

```
Out[21]: Loan_ID          object
Gender          int64
Married         int64
Dependents      int64
Education       int64
Self_Employed  int64
ApplicantIncome int64
CoapplicantIncome float64
LoanAmount      float64
Loan_Amount_Term float64
Credit_History  float64
Property_Area   int64
Loan_Status     int64
LoanAmount_log  float64
TotalIncome     float64
TotalIncome_log float64
dtype: object
```

```

In [22]: #Import models from scikit Learn module:
from sklearn.linear_model import LogisticRegression
from sklearn.cross_validation import KFold #For K-fold cross validation
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier, export_graphviz
from sklearn import metrics

#Generic function for making a classification model and accessing performance:
def classification_model(model, data, predictors, outcome):
    #Fit the model:
    model.fit(data[predictors],data[outcome])

    #Make predictions on training set:
    predictions = model.predict(data[predictors])

    #Print accuracy
    accuracy = metrics.accuracy_score(predictions,data[outcome])
    print ("Accuracy : %s" % "{0:.3%}".format(accuracy))

    #Perform k-fold cross-validation with 5 folds
    kf = KFold(data.shape[0], n_folds=5)
    error = []
    for train, test in kf:
        # Filter training data
        train_predictors = (data[predictors].iloc[train,:])

        # The target we're using to train the algorithm.
        train_target = data[outcome].iloc[train]

        # Training the algorithm using the predictors and target.
        model.fit(train_predictors, train_target)

        #Record error from each cross-validation run
        error.append(model.score(data[predictors].iloc[test,:], data[outcome].iloc[test,:]))

    print ("Cross-Validation Score : %s" % "{0:.3%}".format(np.mean(error)))

    #Fit the model again so that it can be refered outside the function:
    model.fit(data[predictors],data[outcome])

```

C:\anaconda\lib\site-packages\sklearn\cross\_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model\_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.

"This module will be removed in 0.20.", DeprecationWarning)

## Logestic Regression

```
In [24]: outcome_var = 'Loan_Status'
model = LogisticRegression()
predictor_var = ['Credit_History']
classification_model(model, df, predictor_var, outcome_var)
```

Accuracy : 80.945%  
Cross-Validation Score : 80.946%

```
In [25]: #We can try different combination of variables:
predictor_var = ['Credit_History', 'Education', 'Married', 'Self_Employed', 'Property_Area']
classification_model(model, df, predictor_var, outcome_var)
```

Accuracy : 80.945%  
Cross-Validation Score : 80.946%

## Decision Tree

```
In [26]: model = DecisionTreeClassifier()
predictor_var = ['Credit_History', 'Gender', 'Married', 'Education']
classification_model(model, df, predictor_var, outcome_var)
```

Accuracy : 80.945%  
Cross-Validation Score : 80.946%

```
In [27]: #We can try different combination of variables:
predictor_var = ['Credit_History', 'Loan_Amount_Term', 'LoanAmount_log']
classification_model(model, df, predictor_var, outcome_var)
```

Accuracy : 89.414%  
Cross-Validation Score : 68.559%

## Random Forest

```
In [28]: model = RandomForestClassifier(n_estimators=100)
predictor_var = ['Gender', 'Married', 'Dependents', 'Education',
                 'Self_Employed', 'Loan_Amount_Term', 'Credit_History', 'Property_Area',
                 'LoanAmount_log', 'TotalIncome_log']
classification_model(model, df, predictor_var, outcome_var)
```

Accuracy : 100.000%  
Cross-Validation Score : 77.689%

```
In [29]: #Create a series with feature importances:  
featimp = pd.Series(model.feature_importances_, index=predictor_var).sort_values(  
print (featimp)
```

```
Credit_History      0.279387  
TotalIncome_log     0.255215  
LoanAmount_log      0.228393  
Dependents          0.054111  
Property_Area       0.048545  
Loan_Amount_Term    0.041235  
Married             0.026748  
Education           0.024773  
Self_Employed       0.020965  
Gender              0.020629  
dtype: float64
```

```
In [30]: model = RandomForestClassifier(n_estimators=25, min_samples_split=25, max_depth=7  
predictor_var = ['TotalIncome_log', 'LoanAmount_log', 'Credit_History', 'Dependents'  
classification_model(model, df, predictor_var, outcome_var)
```

```
Accuracy : 83.062%  
Cross-Validation Score : 81.109%
```